

Cross-Rover Semantic Segmentation of Mars NavCam Terrain with Domain Adaptation

Cesare Fraccaroli

Computer Vision and Deep Learning
Università degli Studi di Verona (UniVR)


AI4MARS · MSL→MER · ResNet34-U-Net · July 2026

Abstract. This work investigates semantic segmentation of Mars navigation images into four terrain classes (*soil*, *bedrock*, *sand*, *big_rock*) using the AI4MARS dataset. The proposed project pretrains a ResNet34 U-Net on MSL (Curiosity images) and adapts it to MER (Spirit/Opportunity images) through unsupervised domain adaptation (UDA) or a semi-supervised approach. Feature-level UDA is purely unsupervised on MER: it uses labeled MSL data and unlabeled MER images only, with no MER ground-truth labels. We then study how much a small labeled MER set (200 frames) can improve over that pure-UDA setting. The infrequent *big_rock* class stays the most challenging, despite strong transfer performance on the more common classes.

1 Motivation and Rationale

Autonomous Mars rovers like Curiosity (MSL) and Spirit/Opportunity (MER) use navigation cameras for detection, path planning, and terrain segmentation. Terrain segmentation maps each image to terrain types, allowing the rover to trace the optimal path.

The AI4MARS dataset [1] offers expert labels from multiple Mars missions. However, a domain shift exists between missions due to differences in camera geometry, lighting, and terrain types across rovers' cameras. Models trained only on MSL data underperform on MER and do not reach the same performance level as on MSL.



(a) MSL (Curiosity, rock formations)

(b) MER (sand-rich, darker terrain)

Figure 1: Cross-mission appearance shift in NavCam imagery (AI4MARS).

1.1 Why this matters

Manual annotation of Mars images for segmentation is resource intensive, and MER has substantially fewer labeled NavCam frames than MSL. We employ unsupervised domain adaptation (UDA) and semi-supervised learning to leverage unlabeled MER images together with a smaller labeled MER set for semi-supervised training. Pretraining uses the more abundant MSL data. Due to computational constraints, we use only a subset of the AI4MARS dataset.

1.2 Rationale for pipeline design

Instead of training a U-Net from scratch, a ResNet34 encoder is utilized. By implementing this choice we can fine-tune on the specific dataset and avoid the computational burden of training on big datasets, since ResNet34 is already pretrained on ImageNet.

1.3 Problem

In this work, we try to solve the challenge of adapting models from one Mars rover mission (MSL, with labeled data) to another (MER, where most data is unlabeled) using unsupervised domain adaptation (UDA). Both domains share a common set of four terrain classes and one additional null class. Our goal is to train a model that performs well on the target domain (MER) without ever seeing the MER test labels. At the same time, we aim to maintain strong performance on the source domain (MSL).




Figure 2: MER training examples with colored class overlays (image + mask).

2 State of the Art

2.1 ViT models and foundation models

2.1.1 Vision Transformers (ViT)

Vision Transformers (ViT) [5] approach image understanding by dividing each image into fixed-size patches, which are processed as a sequence by transformer encoder blocks. ViT architectures require extensive pre-training and are more resource-demanding than CNNs.

2.1.2 Segment Anything Model (SAM)

SAM [13] is a promptable segmentation foundation model developed by Meta and trained on over 1 billion masks. Given a point, box, or prompt, it produces high-quality object masks with strong zero-shot generalization on natural images. Follow-ups (SAM 2, domain-specific fine-tunes) extend this idea to video and specialized sensors.

2.1.3 Relevance to Mars terrain segmentation

While SAM and ViT-based pipelines represent SOTA segmentation architectures, our project faces important differences and constraints. AI4MARS defines four terrain classes, not prompt-based instance masks as in SAM. Adaptation from labeled MSL to mostly unlabeled MER is the main difficulty. Training budget and labeled MER data are limited, so a ResNet34-U-Net may be more practical than a large transformer or SAM.

CORAL, MMD, and DANN fit CNN encoders by aligning the features in the U-Net bottleneck (`enc4`). Using the same UDA stack on SAM or ViT would need extra work. Foundation models could still serve as frozen encoders or pseudo-label teachers on unlabeled MER frames.

2.2 Unsupervised domain adaptation SOTA

2.2.1 Vision-language models and CLIP

CLIP [7] pretrains joint image and text embeddings, so new categories can be queried at inference with text prompts (class names) instead of extra labeled data. Regarding segmentation and UDA, these models (CLIP, VLMs, etc.) can provide targets for mask generation, pseudo-labels or text-based alignment when both domains share the same class names.

2.2.2 Foundation-model and prompt-based UDA trends

A prevalent trend in modern domain adaptation (UDA) is to use large pretrained models such as CLIP or ViT as fixed backbones, with only minor refinements or fine-tuning of additional layers or tokens for the target domain. When there is scarce labeled data, these models are often used to generate pseudo-labels to facilitate training. There is also increasing interest in employing models like diffusion or SAM to produce coarse masks on new data, which can then be used to train smaller segmentation models alongside a limited set of manually labeled examples.

2.3 Our choices for UDA

The UDA methods below are purely unsupervised on the MER target: MER images enter training without ground-truth masks. We adapt MSL to MER with feature-level methods on shared `enc4` representations.

- **CORAL** [9]: covariance alignment on dense `enc4` feature maps.
- **MMD** [4]: RBF kernel Minimum Mean Discrepancy on dense `enc4` features.
- **DANN** [2]: domain adversarial discriminator that uses gradient reversal with MSL segmentation loss and unlabeled MER batches.
- **ADDA** [10]: two-phase adversarial adaptation: (1) train a domain discriminator on frozen global average pooled features; (2) adapt a joint encoder with MSL segmentation, MER adversarial alignment, and dense CORAL on `enc4` (phase 2).

CORAL, MMD, and DANN combine segmentation loss with a domain adaptation term ($\lambda=0.1$). ADDA instead uses a dedicated alternate schedule (5 discriminator epochs plus 10 adaptation epochs) with adversarial weight ramped by the Ganin schedule and CORAL weight 0.1.

2.3.1 DANN vs ADDA

DANN updates the encoder and discriminator together in one step by using the gradient reversal. ADDA pretrains the discriminator and then alternates between discriminator and encoder updates. Both methods reuse the same ResNet34 encoder.

2.4 Semi-supervised and pseudo-label segmentation

Self-training assigns pseudo-labels to unlabeled pixels that surpass a confidence threshold. We use EMA teachers [11] to stabilize pseudo-labels. We use joint losses that combine labeled and pseudo terms in the same optimization step to try to outperform other approaches. Unlike pure UDA, these methods assume the presence of a small labeled MER set (`mer_labeled_adapt`) and do not require an explicit MSL/MER feature alignment loss.

3 Objectives

3.1 Main objective

Train a segmenter on MSL and adapt it to MER so that MER test mIoU surpasses that of a zero-shot MSL-pretrained baseline while maintaining reasonable MSL test performance. A further goal is to quantify how




Figure 3: Schematic comparison of (a) DANN and (b) ADDA. DANN trains the feature extractor and domain discriminator jointly via gradient reversal; ADDA first trains a discriminator on frozen features, then adapts the target encoder adversarially. Figure adapted from Han et al. [3].

much having labeled MER data helps relative to pure unsupervised domain adaptation.

3.2 Secondary objectives

1. Compare pure unsupervised feature-level UDA methods (CORAL, MMD, DANN, ADDA) that never use MER ground-truth labels.
2. Compare semi-supervised MER variants (pure pseudo-labeling and joint) that use 200 labeled MER frames, and measure their gain over the best pure-UDA result.
3. Analyze class behavior, especially the rare *big_rock* class.

4 Methodology

4.1 Dataset and preprocessing

We use the AI4MARS dataset [1]: RGB images with four terrain classes (*soil*, *bedrock*, *sand*, *big_rock*). Rover body and uncertain pixels (distant objects, sky, background) are excluded from loss and metrics. Images are resized to 512×512 and ImageNet-normalized. On MSL only, auxiliary rover-body and range masks set additional pixels to ignore, avoiding training on the rover or distant regions where manual annotation was deliberately omitted in AI4MARS. We use splits for MSL pretraining and testing, labeled MER adaptation, unlabeled MER UDA, and MER validation and test evaluation.

Table 1: Data splits and roles.

Split	Role	N
msl_train / val / test	Pretrain / test	3000 / 322
mer_labeled_adapt	Semi-supervised anchor	200
mer_uda	Unlabeled MER (UDA + pseudo)	2000
mer_val / mer_test	Select adaptation / final test	200 / 204

4.1.1 Validation protocol

1. Tune on the appropriate validation split.
2. Early-stop on validation mIoU to avoid overfitting.
3. Report test metrics *once* for selected finalists only.

Pretrain hyperparameter search uses `msl_val`.

Table 2: Pretraining hyperparameter sets (winner: low learning rate).

Set	Encoder LR	Decoder LR	Weight decay
Base	10^{-5}	10^{-4}	10^{-4}
Low LR (selected)	5×10^{-6}	5×10^{-5}	10^{-4}
Higher decoder LR	10^{-5}	2×10^{-4}	10^{-4}

4.2 Metrics

We report mean IoU, per-class IoU, and frequency-weighted IoU (wIoU). The problem with unweighted mIoU is that it can mask saturation on frequent classes. Zero-shot pretraining shows a substantial drop from the MSL test to the MER test, illustrating that the cross-mission domain gap is notable.




Figure 4: Domain gap: MSL test vs. MER test wIoU (zero-shot pretrain drops ~ 0.19 on MER).

4.3 Architecture

We couple an ImageNet-pretrained ResNet34 encoder [6] with a U-Net decoder [8] and a segmentation head. Skip connections link encoder stages to matching decoder levels; then, the bottleneck `enc4` contains the deepest features. We split the forward path (encode + decode) so that feature UDA can target `enc4` without breaking segmentation, allowing feature-level UDA to be applied at this bottleneck. **This was the most important implementation choice** in order to make feature-level UDA work.

4.3.1 Encoder-decoder data flow

The input RGB image passes through the ResNet stem and encoder stages, producing skip tensors. The decoder upsamples from `enc4`, concatenating skips at each level and applying convolution blocks at each stage, then a final upsample restores full resolution before the classification head. We apply domain alignment only to `enc4`. The skip tensors retain source fine structure, helping MSL accuracy but limiting how much MER appearance can change in shallow layers.




Figure 5: ResNet34-U-Net architecture: ImageNet-pretrained encoder, skip connections at four scales, bottleneck *enc4* (512×16×16) for feature-level UDA and decoder head producing 512×512 class logits.

4.4 Segmentation loss and pretraining

All training stages minimize cross-entropy on valid pixels. Segmentation loss is the mean cross-entropy over labeled pixels in each batch. Let f_θ denote the segmenter, $\Omega(Y)$ the set of labeled pixels in mask Y , and $\text{CE}(\hat{y}, y)$ the per-pixel cross-entropy. For a batch of images X and labels Y ,

$$\mathcal{L}_{\text{seg}}(X, Y) = \frac{1}{|\Omega(Y)|} \sum_{p \in \Omega(Y)} \text{CE}(f_\theta(X)_p, y_p). \quad (1)$$

Feature UDA methods extract bottleneck *enc4* tensors $\mathbf{z}$, flatten them to per-location features $\phi(\mathbf{z})$ with normalization and apply domain losses on ϕ or on global-average-pooled vectors $h(\mathbf{z}) = \text{GAP}(\mathbf{z})$.

MSL pretraining minimizes $\mathcal{L}_{\text{seg}}$ on the MSL training split. The best checkpoint on MSL validation initializes all adaptation runs, and then hyperparameter search compares learning rates and weight decay. The selected low learning-rate configuration is used to initialize follow-up semi-supervised runs.

4.5 Unsupervised domain adaptation

Here we cover **pure unsupervised** domain adaptation on MER so that the goal is to make *enc4* features **less domain specific** while preserving task accuracy.

4.5.1 Feature-level UDA

From the pretraining checkpoint each iteration pairs **labeled MSL** batches (X_S, Y_S) with **unlabeled MER** images X_T and optimizes

$$\mathcal{L}_{\text{UDA}} = \mathcal{L}_{\text{seg}}(X_S, Y_S) + \lambda \mathcal{L}_{\text{DA}}, \quad \lambda = 0.1, \quad (2)$$

with $\mathcal{L}_{\text{DA}}$ chosen as follows (all on *enc4 feat_map* tensors unless noted).

CORAL [9] aligns **covariance matrices** of source and target. With centered covariance matrices C_S, C_T of $\phi(\mathbf{z}_S)$ and $\phi(\mathbf{z}_T)$,

$$\mathcal{L}_{\text{CORAL}} = \frac{1}{4C} \|C_S - C_T\|_F^2. \quad (3)$$

MMD [4] uses an **RBF-kernel** estimate ($\sigma=1.0$):

$$\mathcal{L}_{\text{MMD}} = \mathbb{E}[k(\phi_s, \phi'_s)] + \mathbb{E}[k(\phi_t, \phi'_t)] - 2\mathbb{E}[k(\phi_s, \phi_t)], \quad k(u, v) = \exp\left(-\frac{\|u-v\|^2}{2\sigma^2}\right). \quad (4)$$

DANN [2] attaches a **domain discriminator** D to enc4 features with **gradient reversal** (GRL) strength augmented by a ramping function with ($\lambda_{\max}=0.1$) that increases over the epochs.

$$\mathcal{L}_{\text{DANN}} = \mathcal{L}_{\text{seg}}(X_S, Y_S) + \text{BCE}(D(\text{GRL}_{\lambda_p}(h_S)), 0) + \text{BCE}(D(\text{GRL}_{\lambda_p}(h_T)), 1), \quad (5)$$

where source labels are 0 and target labels are 1. The discriminator and encoder are updated **jointly** at each step.

Checkpoints for CORAL, MMD, and DANN are selected on MER validation.

4.5.2 ADDA adaptation

ADDA adaptation [10] runs in **two phases** from the same pretrained checkpoint.

Phase 1 (discriminator pre-training). The segmenter is **frozen**, and the discriminator is trained on **GAP(enc4)** features from paired MSL/MER batches, h_S is the global average pooled vector of the enc4 features while h_T is the global average pooled vector of the enc4 features from the target domain (MER). The **global average pooling** is performed in order to map the features to a **512-dimensional** vector instead of flattening a $512 \times 16 \times 16$ tensor:

$$\mathcal{L}_D = \text{BCE}(D(h_S), 0) + \text{BCE}(D(h_T), 1). \quad (6)$$

Phase 2 (encoder adaptation). Each iteration alternates between a discriminator update on detached features and an encoder/decoder step that combines MSL segmentation loss, adversarial loss, and dense CORAL on enc4.

$$\mathcal{L}_{\text{ADDA}} = \lambda_{\text{seg}} \mathcal{L}_{\text{seg}}(X_S, Y_S) + \lambda_{\text{adv}} \mathcal{L}_{\text{adv}} + \lambda_{\text{CORAL}} \mathcal{L}_{\text{CORAL}}, \quad (7)$$

where $\mathcal{L}_{\text{adv}} = \text{BCE}(D(h_T), 0)$ pushes MER features **toward the source domain**, $\mathcal{L}_{\text{CORAL}}$ is Eq. 3 on dense enc4 maps, $\lambda_{\text{seg}}=1$, λ_{adv} is ramped (Ganin schedule, peak 0.1) and $\lambda_{\text{CORAL}}=0.1$. The dense CORAL term is **our addition** to the classic ADDA approach, which uses the adversarial loss alone. We keep it because it **stabilizes** the encoder update on our small batches. The selected recipe uses 5 discriminator epochs, 10 adaptation epochs and batch size 32.

4.6 Semi-supervised adaptation on MER

After measuring what unsupervised adaptation can achieve without MER labels we ask how much a **small labeled MER set** adds on top of it. We adapt the segmenter on the MER target domain using the 200 labeled MER frames (`mer_labeled_adapt`) and the 2000 unlabeled MER frames (`mer_uda`). All variants are initialized from the **MSL pretraining checkpoint** (zero-shot).

4.6.1 Pseudo-labeling

Pure pseudo-labeling is self-training without MER ground-truth anchors. On unlabeled MER images X_T , the model (or EMA teacher) produces per-pixel class predictions $\hat{y}_p = \arg \max_c f_{\theta}(X_T)_{p,c}$ and confidence $c_p = \max_c \text{softmax}(f_{\theta}(X_T))_{p,c}$. Pixels with $c_p < \tau$ are ignored and batches with too few confident pixels are skipped. Self-training minimizes

$$\mathcal{L}_{\text{pseudo}} = \mathcal{L}_{\text{seg}}(X_T, \hat{Y}_T), \quad (8)$$

on pseudo-labels. When an EMA teacher θ' is used, teacher weights are updated after each step with exponential moving average decay towards the student.

Without labeled MER anchoring, confident errors can reinforce themselves and the method **collapses** on MER test (Table 4).

4.6.2 Joint semi-supervision

We use a combination of labeled and pseudo terms in a **single optimization step** (inspired by the Mean Teacher approach). The joint variant combines labeled and pseudo terms in a single optimization step:

$$\mathcal{L}_{\text{joint}} = \mathcal{L}_{\text{seg}}(X_T^{\text{lab}}, Y_T^{\text{lab}}) + \lambda_u \mathcal{L}_{\text{pseudo}}(X_T^{\text{uda}}), \quad (9)$$

with an **EMA teacher**, **two labeled-only warm-up epochs**, and the same confidence filtering as Eq. 8. Pipeline (one training iteration after warm-up):

1. Sample a labeled MER batch ($X_T^{\text{lab}}, Y_T^{\text{lab}}$) and an unlabeled MER batch X_T^{uda} .
2. Compute $\mathcal{L}_{\text{seg}}$ on the labeled batch with true MER masks.
3. Generate pseudo-labels on X_T^{uda} with the student (or EMA teacher); drop pixels with confidence $< \tau$.
4. Compute $\mathcal{L}_{\text{pseudo}}$ on the remaining confident pixels.
5. Form $\mathcal{L}_{\text{joint}} = \mathcal{L}_{\text{seg}} + \lambda_u \mathcal{L}_{\text{pseudo}}$ and run one backward update.
6. If an EMA teacher is used, update it toward the student after the step.

5 Experiments and Results

5.1 Experimental setup

All runs use AdamW. Full training first was tried on a rtx 2060 local eGPU but was latter discarded due to memory and time constraints and used a cloud-hosted RTX 4090 yielding a 20x improvement in training time.

Table 3: Adaptation hyperparameters per method.

Method	Batch	Epochs	LR	Loss weights
CORAL / MMD / DANN	16	10	0.1×	$\lambda_{\text{DA}}=0.1$
ADDA	32	5+10	own	$\lambda_{\text{seg}}=1, \lambda_{\text{adv}}=\lambda_{\text{CORAL}}=0.1$
Pseudo-label	16	10	0.1×	pseudo only
Joint semi-supervision	16	2+10	0.05×	$\lambda_u=0.5$ (ramp 4 ep.)

5.2 Implementation overview

All experiments share a single PyTorch pipeline built around the same **ResNet34-U-Net** segmenter, and every method reuses it as a baseline, differing only in the training objective (MSL pretraining, feature-level alignment on enc4, pseudo-label self-training, or joint semi-supervision). Checkpoints store only model weights; the best epoch is chosen, and metrics are logged separately for analysis.

5.3 Pretraining

MSL pretraining on the training split reaches strong validation mIoU; the selected checkpoint is evaluated **zero-shot** on MER before any adaptation (Table 4, zero-shot row).

5.4 Main pipeline

Table 4 summarizes test mIoU and wIoU after one validation-selected evaluation per method.



On the main pipeline, **zero-shot** pretraining establishes a baseline on the MER test. Feature-level UDA methods (CORAL, MMD, DANN, ADDA) **improve** over zero-shot, with **ADDA** performing best among pure UDA approaches on weighted IoU. Among UDA-only methods, **ADDA** also attains the highest MER mIoU in this table. Pure pseudo-labeling **collapses** without labeled MER anchoring. **Joint semi-supervision** (reported separately in Table 5, since it uses MER labels) reaches the **best overall MER test mIoU** by combining labeled and unlabeled MER data in one optimization step. MMD and DANN also show higher MSL test than zero-shot because UDA continues MSL supervised training during adaptation while selecting on MER validation.

5.5 Semi-supervised and pseudo-labeling

From the pretraining checkpoint, we tried pseudo-labeling with an EMA teacher and joint semi-supervision. Pseudo-labeling also employs an EMA teacher but does not use the 200 true MER labels. The student is trained on the teacher’s pseudo-labels. In **joint semi-supervision**, every update also uses the **200 true MER labels** so the model does not rely only on pseudo-labels. This joint setup gives the **highest MER test mIoU** among the methods we tried. Per-class gains are largest for the rare *big_rock* class relative to zero-shot pretraining, though absolute *big_rock* IoU remains **low**. Table 5 lists follow-up runs; Table 7 reports per-class MER test IoU for joint semi-supervision vs. zero-shot pretrain. Relative to Table 4, Table 5 focuses on the SSL

Table 4: Main pipeline: test mIoU and weighted IoU (wIoU).

Method	MSL test			
	mIoU	wIoU	mIoU	wIoU
Zero-shot (pretrain)	0.764	0.958	0.590	0.769
CORAL	0.797	0.955	0.623	0.810
MMD	0.772	0.947	0.639	0.821
DANN	0.828	0.957	0.637	0.820
ADDA	0.857	0.944	0.647	0.846
Pseudo-label	0.136	0.273	0.127	0.203

**Figure 6:** Zero-shot U-Net on MER test (MSL-pretrained, no adaptation): original NavCam frame, ground-truth evaluation mask and predicted mask.**Figure 7:** Feature-level UDA only on MER test: mIoU and wIoU (zero-shot, CORAL, MMD, DANN, ADDA).

family and shows that **EMA alone is not enough**: joint labeled+pseudo optimization is required for the best MER mIoU.

Table 5: Semi-supervised follow-up: MER test mIoU and wIoU.

Method	mIoU	wIoU
Zero-shot (pretrain), reference	0.590	0.769
ADDA (best pure UDA), reference	0.647	0.846
Pseudo-label + EMA	0.618	0.803
Pseudo-label + EMA (strict)	0.590	0.772
Joint semi-supervision	0.658	0.844




Figure 8: Semi-supervised MER adaptation only: mIoU and wIoU (zero-shot, pure pseudo-label, Pseudo+EMA, joint).

Frequent classes achieve strong IoU even in the zero-shot setting; performance improvements are primarily observed for the *big_rock* class.

5.6 Per-class analysis

Feature-level UDA improves **frequent-class** IoU on the MER test, with method-dependent trade-offs between macro mIoU and weighted IoU. Table 6 breaks down CORAL, MMD, DANN and ADDA on MER test with per-class IoU, mIoU and wIoU. Comparing zero-shot pretraining with joint semi-supervision shows gains across all four terrain classes, with the **largest relative improvement** on *big_rock*. Table 7 compares zero-shot pretrain and joint semi-supervision on MER test. Improvements for the *big_rock* class remain **limited** in comparison to the frequent classes.

Table 6: Feature-level UDA on MER test: per-class IoU, mIoU and wIoU.

Method	soil	bedrock	sand	big_rock	mIoU	wIoU
CORAL	0.817	0.827	0.812	0.036	0.623	0.810
MMD	0.823	0.823	0.844	0.066	0.639	0.821
DANN	0.828	0.789	0.866	0.066	0.637	0.820
ADDA	0.865	0.798	0.893	0.033	0.647	0.846

5.7 Discussion of results

Pseudo-labeling with an EMA teacher reaches moderate mIoU but remains **inferior** to joint semi-supervision. CORAL underperforms relative to MMD and DANN on MER mIoU, while **ADDA** attains the highest wIoU among UDA methods.




Figure 9: Qualitative MER test predictions after adaptation (image, GT and prediction). Top: DANN feature UDA. Bottom: joint semi-supervision.

Table 7: Per-class IoU on MER test (zero-shot pretrain vs joint semi-supervision).

Class	Zero-shot (pretrain)	Joint semi-supervision
soil	0.788	0.855
bedrock	0.752	0.787
sand	0.784	0.907
big_rock	0.036	0.084
mIoU	0.590	0.658
wIoU	0.769	0.844




Figure 10: Per-class IoU on MER test across methods (grouped bars).

ADDA is the strongest pure-UDA method on MER wIoU (**0.846**), while **joint semi-supervision** reaches higher MER mIoU (**0.658** vs. 0.647 for **ADDA**) when 200 labeled MER frames are available.

The aggregate gain from those labels is **small**: +0.011 mIoU and −0.002 wIoU, i.e. joint semi-supervision is marginally better on the macro average and marginally worse on the frequency-weighted one. The difference is concentrated almost entirely in the **rare class**: joint semi-supervision reaches *big_rock* IoU **0.084**, whereas **ADDA** reaches 0.033, which is below the 0.036 of the unadapted model. In other words, adversarial feature alignment buys its high wIoU on the **frequent classes** and does not help *big_rock*, while the 200 MER labels act mainly as **rare-class supervision** rather than as a general accuracy boost.

6 Conclusions

We addressed cross-rover Mars NavCam segmentation with a ResNet34-U-Net pretrained on MSL and adapted to MER using feature-level UDA and semi-supervised joint training on AI4MARS. The two adaptation families were studied separately: **pure UDA** without MER labels, and **semi-supervised** MER adaptation with 200 labeled MER images to quantify the gain of labels over unsupervised adaptation. **Zero-shot** transfer from pretraining establishes a usable MER baseline. Feature UDA improves MER test mIoU over zero-shot, with **ADDA** also yielding the highest UDA wIoU. **Joint semi-supervision** further improves MER mIoU over the best pure-UDA result by using the labeled MER set together with unlabeled frames, but only by +0.011 mIoU, and with a marginally lower wIoU; the 200 MER labels therefore add little on the aggregate metrics and pay off mainly on the **rare class**. The *big_rock* class remains **difficult** because it covers only a tiny fraction of pixels; mIoU gains partly reflect rare-class recovery under joint semi-supervision.

Joint semi-supervision leads on macro mIoU over zero-shot, while **ADDA** leads on wIoU; **joint semi-supervision** is the only method that pushes *big_rock* IoU beyond 0.08, although **MMD** and **DANN** already nearly double it (0.036 to 0.066) with **no MER labels** at all, among the methods we consider deployable (zero-shot pretrain plus feature UDA and joint semi-supervision). Figure 12 consolidates these comparisons on MER test.

6.1 Future work

1. SAM or CLIP encoders could provide pseudo-masks or semantic targets for unlabeled MER frames, complementing the current EMA self-training loop without replacing the lightweight U-Net deployment path.
2. *m2020_val* (Perseverance rover) is not considered in our protocol; evaluating if MSL→MER models transfer to M2020 without retraining would test cross-mission robustness beyond the Spirit/Opportunity target.
3. A natural next step is to **combine pure UDA and joint semi-supervision** (e.g. by initializing joint training from the best UDA checkpoint instead of from MSL pretrain only).




Figure 11: Row-normalized confusion matrices on MER test (class recall).




Figure 12: Summary comparison on MER test: mIoU vs. wIoU (A), per-class IoU heatmap (B), *big_rock* IoU (C) and *big_rock* recall from confusion matrices (D).

References

- [1] A. Azari, B. Chen, A. Daftry, et al. AI4MARS: A dataset for terrain-aware autonomous driving on Mars. In *NeurIPS Datasets and Benchmarks Track*, 2021.
- [2] Y. Ganin, E. Ustinova, H. Ajakan, et al. Domain-adversarial training of neural networks. *JMLR*, 17(1):2096-2030, 2016.
- [3] C. Han, H. Choo and J. Jeong. Bidirectional-feature-learning-based adversarial domain adaptation with generative network. *Applied Sciences*, 13(21):11825, 2023.
- [4] A. Gretton, K. M. Borgwardt, M. J. Rasch, B. Schölkopf and A. Smola. A kernel two-sample test. *JMLR*, 13:723-773, 2012.
- [5] A. Dosovitskiy, L. Beyer, A. Kolesnikov, D. Weissenborn, X. Zhai, T. Unterthiner, M. Dehghani, M. Minderer, G. Heigold, S. Gelly, J. Uszkoreit and N. Houlsby. An image is worth 16x16 words: Transformers for image recognition at scale. In *ICLR*, 2021.
- [6] K. He, X. Zhang, S. Ren and J. Sun. Deep residual learning for image recognition. In *CVPR*, 2016.
- [7] A. Radford, J. W. Kim, C. Hallacy, A. Ramesh, G. Goh, S. Agarwal, G. Sastry, A. Askell, P. Mishkin, J. Clark, G. Krueger and I. Sutskever. Learning transferable visual models from natural language supervision. In *ICML*, 2021.
- [8] O. Ronneberger, P. Fischer and T. Brox. U-Net: Convolutional networks for biomedical image segmentation. In *MICCAI*, 2015.
- [9] B. Sun and K. Saenko. Deep CORAL: Correlation alignment for deep domain adaptation. In *ECCV Workshops*, 2016.
- [10] E. Tzeng, J. Hoffman, K. Saenko and T. Darrell. Adversarial discriminative domain adaptation. In *CVPR*, 2017.
- [11] A. Tarvainen and H. Valpola. Mean teachers are better role models. In *NeurIPS*, 2017.
- [12] Y.-H. Tsai, W.-C. Hung, S. Schultter, M. Yang, M.-H. Yang and M. Chandraker. Learning to adapt structured output space for semantic segmentation. In *CVPR*, 2018.
- [13] A. Kirillov, E. Mintun, N. Ravi, H. Mao, C. Rolland, L. Gustafson, T. Xiao, S. Whitehead, A. C. Berg, W.-Y. Lo, P. Dollar and R. Girshick. Segment Anything. In *ICCV*, 2023.